

Pintos project 3

Project 2-2 is done and before you start project 3 ..

- Make a new branch named “project3” and do your work at that branch.
 - git branch project3
 - git checkout project3
- Without finishing 2-2, you can't do 3. Please finish 2-2 first.

• **Make commits frequently!**

• **Due**

- **6/11(Wednesday)**

Pre-work

- Read pintos manual.
 - Project 3: Virtual Memory
 - https://web.stanford.edu/class/cs140/projects/pintos/pintos_4.html#SEC53
- There is no files in the vm directory.
 - You should write your own files in the vm directory from scratch.
(*frame.c, frame.h, page.c, page.h, swap.c, swap.h*)
 - Also, add your own files (*src/Makefile.build*)
 - `$ vi Makefile.build`

```
# No virtual memory code yet.  
#vm_SRC = vm/file.c      # Some file.  
vm_SRC += vm/page.c  
vm_SRC += vm/frame.c  
vm_SRC += vm/swap.c
```

Goal

- **Implement Virtual Memory**

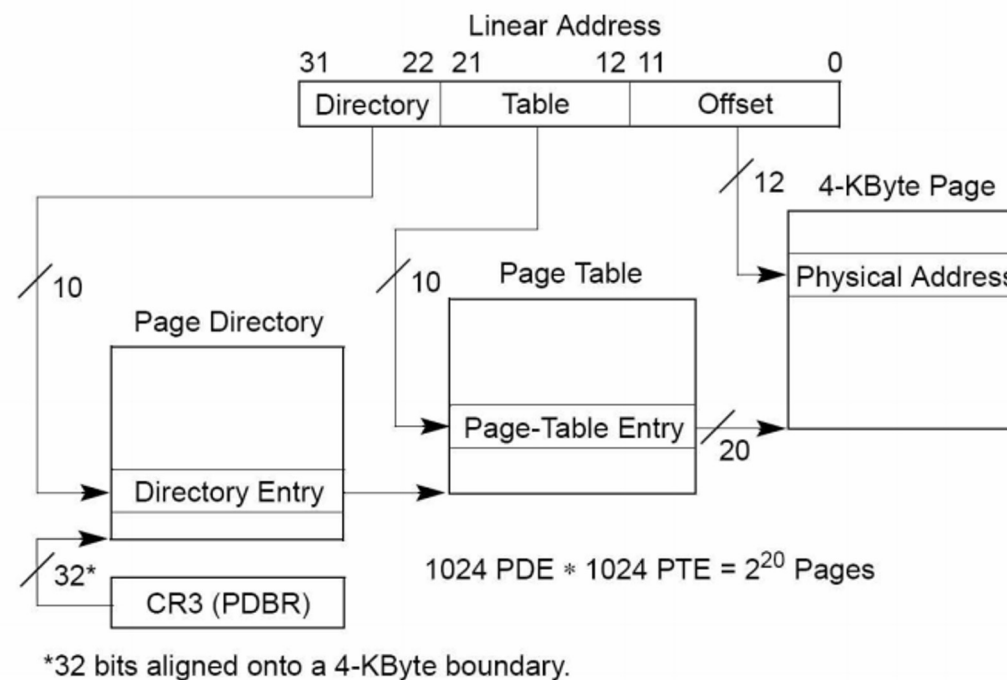
- Page Table Management
 - Page fault handler
 - Eviction policy
- Swap Disk Management
 - Evicted pages
- Stack Growth
- Memory Mapped Files

Suggested Order of Implementation

- Follow “4.2 Suggested Order of Implementation” in pintos manual
 - Implement Frame Table.
 - See section “4.1.5 Managing the Frame Table”
 - Modify *userprog/process.c* to use your frame table allocator.
 - After this step, your kernel still pass all the project 2 test cases.
 - Do not implement swapping yet
 - Supplemental page table and page fault handler
 - See section “4.1.4 Managing Supplemental Page Table”
 - Modify *userprog/process.c* to record the necessary information in supplemental page table when loading an executable and setting up its stack
 - Implement loading of code and data segments in the page fault handler
- Implement Swapping, Stack Growth, Mapped Files

Page Table Management

- Paging in the x86 architecture



Page Table Management

- Current Pintos Implementation
 - Use paging
 - Page size: 4KB
 - Each process has its own page tables
- The page directory is allocated when the process is created
(`pagedir_create()` - *userprog/pagedir.c*)
- Thread → `pagedir` points to the page directory
(`load()` - *userprog/process.c*)
- The (secondary) page tables are dynamically created if necessary
(`lookup_page()` - *userprog/pagedir.c*)
- For kernel region, processes have the same mapping
(`PHYS_BASE ~ 0xffffffff`)

Page Table Management

- Current Pintos Implementation
 - No demand paging
 - When a process is created, all the contents of code and data segments are read into the physical memory
(`load_segment()` – *userprog/process.c*)
 - Fixed stack size
 - Only one stack page is allocated to each process
(`setup_stack()` – *userprog/process.c*)
 - No page faults in the user mode
 - Everything needed by each process is in the physical memory

Lazy Loading (Demand Paging)

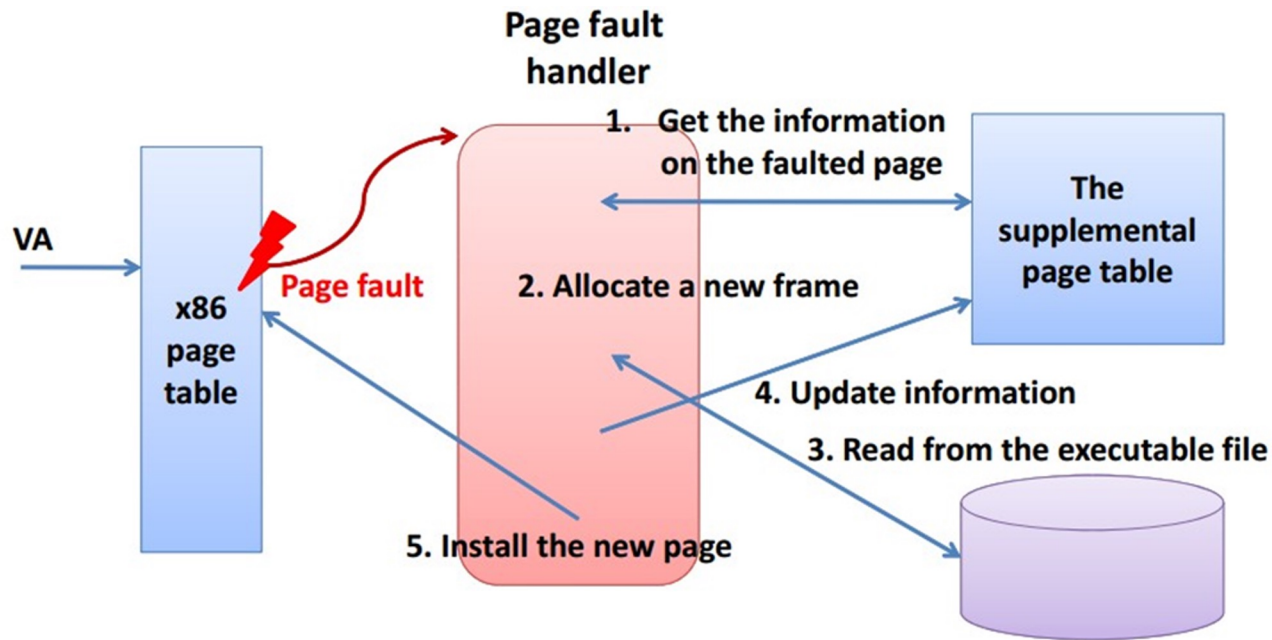
- Loading code/data from the executable file
 - (`load_segment()` – *userprog/process.c*)
- A process do not need all the pages immediately
- Use the executable file as the backing store
 - Only when a page is needed at run time, load the corresponding code/data page into the physical memory
 - Loaded pages will have valid PTEs
- Supplemental page table
 - The page table with additional data about each page
 - On a page fault, find out what data should be there for the faulted virtual page
 - On a process termination, decide what resources to free
 - Any data structure for the supplemental page table
(refer to *lib/kernel/hash.c*)

Lazy Loading (Demand Paging)

- Page Fault Handler (section 4.1.4)
 - Implement by modifying `page_fault()` in *userprog/exception.c*
 - Locate the page that faulted in the supplemental page table.
 - If the memory reference is valid
 - Obtain a frame to store the page
 - Fetch the data into the frame
 - Point the page table entry for the faulting virtual address to the physical page
(use the functions in *userprog/pagedir.c*)
 - Any invalid access terminates the process and thereby frees all of its resources


Lazy Loading (Demand Paging)

- Page Fault Procedure



Swapping

- Swap disk
 - Used when the amount of physical memory is full
 - If the system needs more memory resources and the RAM is full,
 - Inactive pages in memory are moved to the swap space



Swapping

- Create Swap disk
 - Use the following command to create an 4 MB swap disk in the *vm/build* directory
 - `$ pintos -mkdisk swap.dsk 4`
- Accessing swap disk
 - The swap disk is automatically attached as hd1:1 when you run Pintos
 - Use the disk interface in *devices/disk.h*
 - `struct disk *disk_get (int chan_no, int dev_no);`
 - `disk_sector_t disk_size (struct disk *);`
 - `void disk_read (struct disk *, disk_sector_t, void *);`
 - `void disk_write (struct disk *, disk_sector_t, const void *);`

Swapping

- A swap disk consists of swap slots
 - A swap slot is a continuous, page-size region of disk space on the swap disk



- Managing swap slots
 - Pick an unused swap slot for evicting a page
 - Free a swap slot when its page is read back or the process is terminated
- Swap table
 - The swap table tracks in-use and free swap slots
 - [*lib/kernel/bitmap.c*](#) will be useful

Swapping

- Page replacement policy
 - Implement global page replacement algorithm that approximates LRU
 - The “second chance” or “clock” algorithm
- Get/Clear Accessed and Dirty bits in the PTE
 - In *userprog/pagedir.c*,
 - `pagedir_is_dirty()`, `pagedir_set_dirty()`
 - `pagedir_is_accessed()`, `pagedir_set_accessed()`

Swapping


- Frame Table
 - Allows efficient implementation of eviction policy
 - Use the frame table while you choose a victim page to evict when no frames are free
 - One entry for each frame that contains a user page
- Frame Allocation
 - `pallocc_get_page()` and `pallocc_free_page()` in *threads/pallocal.c*
 - If there are free frames in the user pool, allocate one by calling `pallocc_get_page()`
 - If none is free
 - Choose a victim page using your page replacement policy
 - Remove references to the frame from any page table that refers to it
 - If the frame is modified, write the page to the file system or to the swap disk
 - Return the frame

Stack Growth

- Growing the stack segment
 - Grow the stack dynamically, as page faults occur
 - Identify stack accesses by looking at the `esp` of the struct `intr_frame`
 - When the page fault occurred in the user mode
 - Use `(struct intr_frame *) f->esp`
 - When the page fault occurred in the kernel mode
 - Save `esp` into struct thread on the initial transition from user to kernel mode
- Set an absolute limit on stack size, 8 MB

Memory Mapped Files

- Memory mapped file
 - Map files into process address space
 - On page fault, load file data into page
 - When swapping, write file pages back to disk



Memory Mapped Files

- Implement `mmap()` System calls (In *userprog/syscall.c*)
 - `mapid_t mmap` (`int fd, void *addr`)
 - Need to connect mapids to actual files
 - `mmap()` fails if
 - `fd` is 0 or 1
 - The file has a length of zero bytes
 - `addr` is 0
 - `addr` is not page-aligned
 - The range of pages mapped overlaps any existing set of mapped pages
 - `void munmap` (`mapid_t mapping`)
 - All mapped files should be unmapped when a process exits

Grading

- Code (pintos grade) : 90%
 - **Pass all test cases**
- Design document : 10%
 - Answer questions:
 - Please answer as many questions as you can even if you can't implement

About questions

- We will not answer the question about detailed implementation and code, etc.
- Please ask questions only related to environment and high-level design of project.
- Please read pintos manual carefully.

Thank you